

# Hackapizza 2025

Team - LSF

# Summary

- Problem Description
- Solution Description
  - Design Principles
  - High Level Architecture
  - Offline Phase: Knowledge Base Creation
  - Online Phase: Knowledge Base Querying
- Results & Improvements



# Problem Description

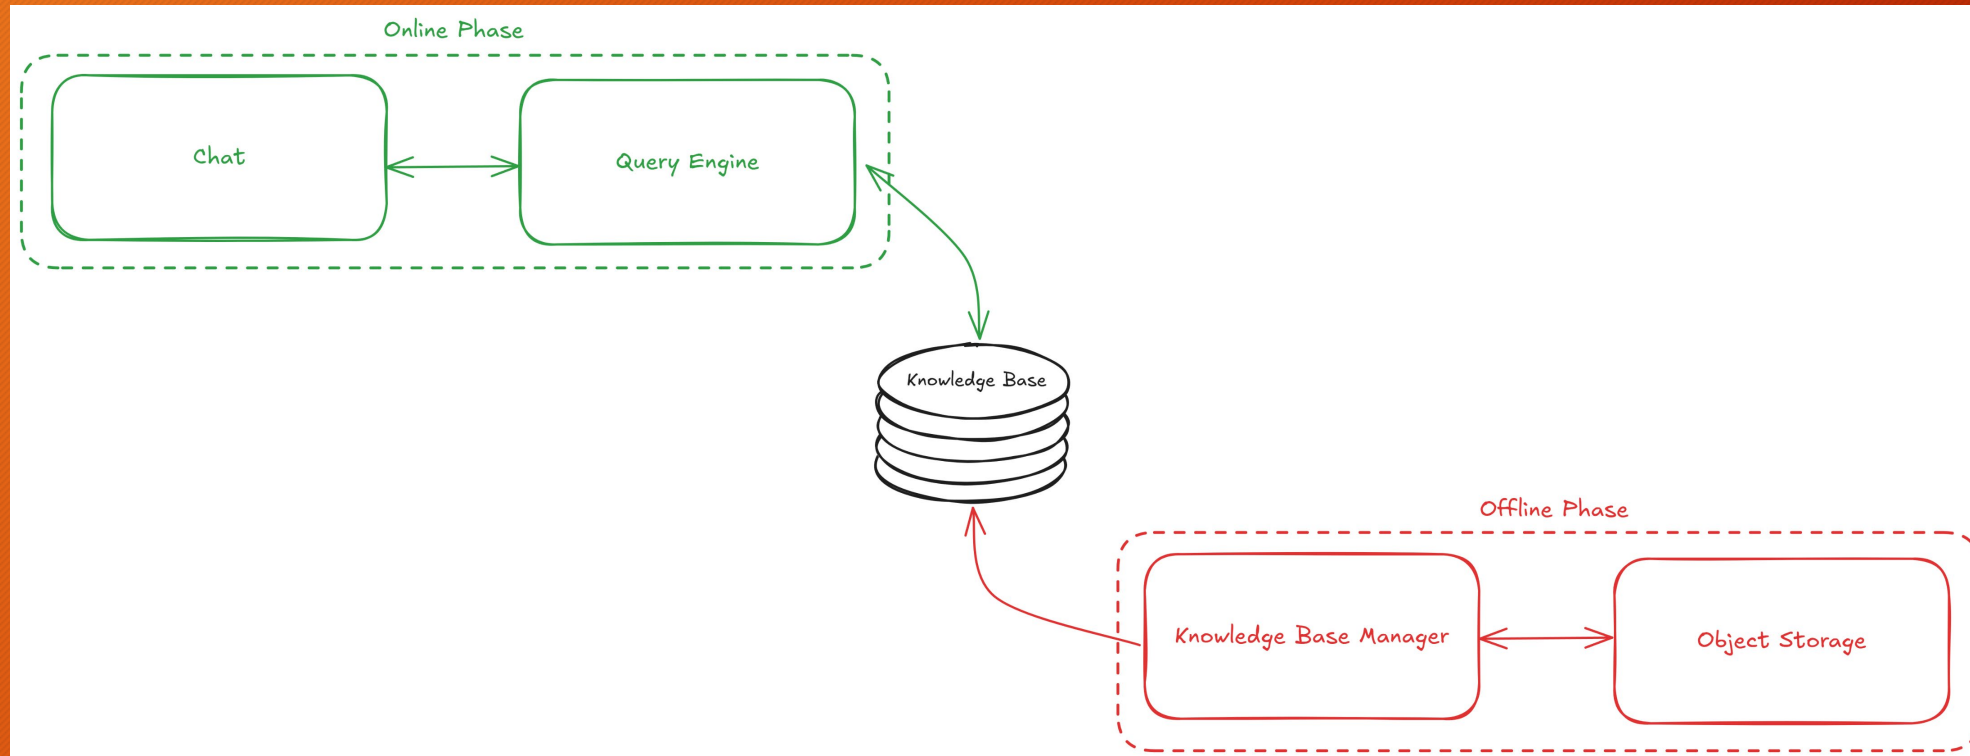
- **Goal:** develop a “intergalactic food advisor” that can provide the user with a list of suitable dishes based on a question in natural language, which implicitly contains some constraints.
- **Basic Setting:** in the galaxy there are 30 different *restaurants* scattered around 10 different *planets*, each with its *chef* and its unique set of *dishes* (i.e., menu). Every *chef* has their own *cooking* licenses, and every *dish* is prepared with specific *ingredients* and *cooking techniques*.
- **Additional Info:** every *cooking technique* requires a set of *licenses* to be performed, some *ingredients* that contain weird *substances* cannot exceed a certain thresholds, and there are three different *orders* to take care of.

# Solution Description - Design Principles

- *Make the System as Predictable and Controllable as Possible.*
  - Combine LLMs with Rule-Based Parsing.
  - Map Questions to Predefined Sequences of Actions.
- *Make the System Deployable On-Premises with Limited Resources.*
  - Use “Small” Open Weights Large Language Models (LLMs).
- *Focus on Data Quality.*
  - Define a Precise Data Model that Suits the Application.



# Solution Description - High Level Architecture



# Solution Description - Offline Phase (1/2)

- **Goal:** creating the Knowledge Base.
- **Description:** the *Knowledge Base Manager* (i) reads *menus* and supporting documents (e.g., *cooking manual*, *intergalactic code of conduct*, *dish mappings*, *planets distances*, ...) from an **object storage**, (ii) generates for each *dish* a JSON descriptor that represents it, (iii) and saves it in a **document database**.
- **Notes:** in this exemplified setting both the **object storage** and the **document database** are folders in our repo, but the switch to actual tools such as *MinIO* and *MongoDB*, for instance, is simple.

# Solution Description - Offline Phase (2/2)

- **Data Model:** for each *dish*, we create a descriptor with the following information.
  - **Restaurant**
    - Name
    - Planet
  - **Chef**
    - Name
    - Licenses (Name, Code, Level)
  - **Dish**
    - Code
    - Name
    - Ingredients
    - Techniques (Name, Category, Subcategory)
  - **Orders Info**
    - *Andromeda* Info
    - *Armonisti* Info
    - *Naturalisti* Info



# Solution Description - Online Phase (1/2)

- **Goal:** querying the Knowledge Base.
- **Description:** the *Query Engine* (i) receives a *question* from the user, (ii) “understands” it, (iii) and transforms it in a sequence of filters to apply on the *Knowledge Base* so to extract the *dishes* that respect the question’s constraints.
- **Notes:** in this exemplified setting we made some assumptions about the types of questions the system can handle. Nevertheless, with some work, one can make the whole process more general and robust.



# Solution Description - Online Phase (2/2)

- **Question Understanding:** each question is processed to produce a descriptor that strictly parallel that of a dish by applying the following steps.
  - Extract *ingredients* and *techniques* (LLM-Based).
  - Extract *restaurant* and *planet* (Rule-Based).
  - Extract *licenses* (LLM-Based).
  - Extract *orders* (Rule-Based).
- **Question Execution:** the content of the *Knowledge Base* is recursively filtered based on the *question* descriptor extracted in the previous step to produce a list of suitable *dishes*.

# Solution Description - Results & Improvements

- The results are promising, considering the limited resources we relied upon. In particular, we obtained a **72.5** score with our best submission using the approach described in this presentation.
- Here's some potential improvements:
  - Refine the question understanding and execution mechanisms to produce more accurate questions representations and more reliable sequences of action.
  - Use a more powerful model to answer the question by passing it to a more powerful model with the whole *Knowledge Base* and leveraging the Data Model.